

Customer churn rate analysis

Objective : To understand the churn rate of customers based on other factors

Algorithms used: Used decision tree initially and checked the accuracy and other factors , however the pictorial representation of tree was not very informative due to consideration of multiple features. Therefore proceeded to use a random forest and logistic regression to understand the important features.

Steps involved

1)importing all the libraries

```
#using decision trees to analyse customer churn rate  
#importing the necessary libraries  
import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt
```

2)importing data and reading data for first 5 rows

```
[2]: #reading the data
data = pd.read_csv("WA_Fn-UseC_-Telco-Customer-Churn.csv")

[3]: data.head()
```

| | customerID | gender | SeniorCitizen | Partner | Dependents | tenure | PhoneService | MultipleLines | InternetService | OnlineSecurity | ... | DeviceProtection | TechSupport | S |
|---|------------|--------|---------------|---------|------------|--------|--------------|------------------|-----------------|----------------|-----|------------------|-------------|---|
| 0 | 7590-VHVEG | Female | 0 | Yes | No | 1 | No | No phone service | DSL | No | ... | No | No | |
| 1 | 5575-GNVDE | Male | 0 | No | No | 34 | Yes | No | DSL | Yes | ... | Yes | No | |
| 2 | 3668-QPYBK | Male | 0 | No | No | 2 | Yes | No | DSL | Yes | ... | No | No | |
| 3 | 7795-CFOCW | Male | 0 | No | No | 45 | No | No phone service | DSL | Yes | ... | Yes | Yes | |
| 4 | 9237-HQITU | Female | 0 | No | No | 2 | Yes | No | Fiber optic | No | ... | No | No | |

3)checking for null values as well the statistical parameters of the data , using functions info and describe

```
[4]: data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
 #   Column                Non-Null Count  Dtype  
---  --
 0   customerID            7043 non-null   object  
 1   gender                 7043 non-null   object  
 2   SeniorCitizen          7043 non-null   int64   
 3   Partner                7043 non-null   object  
 4   Dependents             7043 non-null   object  
 5   tenure                 7043 non-null   int64   
 6   PhoneService           7043 non-null   object  
 7   MultipleLines          7043 non-null   object  
 8   InternetService        7043 non-null   object  
 9   OnlineSecurity         7043 non-null   object  
10   OnlineBackup           7043 non-null   object  
11   DeviceProtection       7043 non-null   object  
12   TechSupport            7043 non-null   object  
13   StreamingTV            7043 non-null   object  
14   StreamingMovies        7043 non-null   object  
15   Contract               7043 non-null   object  
16   PaperlessBilling       7043 non-null   object  
17   PaymentMethod          7043 non-null   object  
18   MonthlyCharges         7043 non-null   float64  
19   TotalCharges           7043 non-null   object  
20   Churn                  7043 non-null   object  
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB
```

```
[5]: data.describe()
```

```
[5]:
```

| | SeniorCitizen | tenure | MonthlyCharges |
|-------|---------------|-------------|----------------|
| count | 7043.000000 | 7043.000000 | 7043.000000 |
| mean | 0.162147 | 32.371149 | 64.761692 |
| std | 0.368612 | 24.559481 | 30.090047 |
| min | 0.000000 | 0.000000 | 18.250000 |
| 25% | 0.000000 | 9.000000 | 35.500000 |
| 50% | 0.000000 | 29.000000 | 70.350000 |
| 75% | 0.000000 | 55.000000 | 89.850000 |
| max | 1.000000 | 72.000000 | 118.750000 |

statistical values like the standard deviation and inter quartile ranges can help us understand the spread of the data and hence understand the range and outliers.

4)checking for null values and removing all the NA values

```
[84]: print(data.isnull().sum())
```

```
customerID    0
gender        0
SeniorCitizen 0
Partner       0
Dependents    0
tenure        0
PhoneService  0
MultipleLines 0
InternetService 0
OnlineSecurity 0
OnlineBackup  0
DeviceProtection 0
TechSupport   0
StreamingTV   0
StreamingMovies 0
Contract      0
PaperlessBilling 0
PaymentMethod 0
MonthlyCharges 0
TotalCharges  0
Churn         0
dtype: int64
```

```
[93]: data.dropna().head()
```

```
[93]:
```

| | gender | SeniorCitizen | Partner | Dependents | tenure | PhoneService | MultipleLines | InternetService | OnlineSecurity | OnlineBackup | DeviceProtection | TechSupport | StreamingTV |
|---|--------|---------------|---------|------------|--------|--------------|---------------|-----------------|----------------|--------------|------------------|-------------|-------------|
| 0 | 0 | 0 | 1 | 0 | 1 | 0 | 1 | 0 | 0 | 2 | 0 | 0 | 0 |
| 1 | 1 | 0 | 0 | 0 | 34 | 1 | 0 | 0 | 2 | 0 | 2 | 0 | 0 |
| 2 | 1 | 0 | 0 | 0 | 2 | 1 | 0 | 0 | 2 | 2 | 0 | 0 | 0 |
| 3 | 1 | 0 | 0 | 0 | 45 | 0 | 1 | 0 | 2 | 0 | 2 | 2 | 0 |
| 4 | 0 | 0 | 0 | 0 | 2 | 1 | 0 | 1 | 0 | 0 | 0 | 0 | 0 |

5)label encoder gives unique numerical values to different strings in the data set

```
#Label encoder to convert non numerical data to numerical data
```

```
from sklearn.preprocessing import LabelEncoder
```

```
label_encoders = {}
```

```
#selecting the columns with string datas instead of numerical data
```

```
categorical_columns = data.select_dtypes(include =['object']).columns
```

```
for columns in categorical_columns:
```

```
    if columns != 'customerID':
```

```
        label_encoders[columns] = LabelEncoder()
```

```
        data[columns] = label_encoders[columns].fit_transform(data[columns])
```

6) dropping the unnecessary columns

```
# Drop customerID as it's not a feature
data = data.drop(['customerID'], axis=1)

# Define features and target
X = data.drop('Churn', axis=1) # Features
y = data['Churn'] # Target
```

All this comes under cleaning and processing the data , the next step is to train the algorithm with data and so the data needs to be split into two sets , test set and training set

7)Importing the train_test_split function from sklearn.model_selection library

```
#dividing the given data into training and test set , test size = 0.2 implies 20 percent of the data will be used for testing
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

8)training the decision tree on the given data and checking the parameters to understand how effectively the algorithm works

```
01: #Importing the decision tree classifier
from sklearn.tree import DecisionTreeClassifier

decision_tree = DecisionTreeClassifier(random_state=42)
decision_tree.fit(X_train, y_train)

02: DecisionTreeClassifier
DecisionTreeClassifier(random_state=42)

03: from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

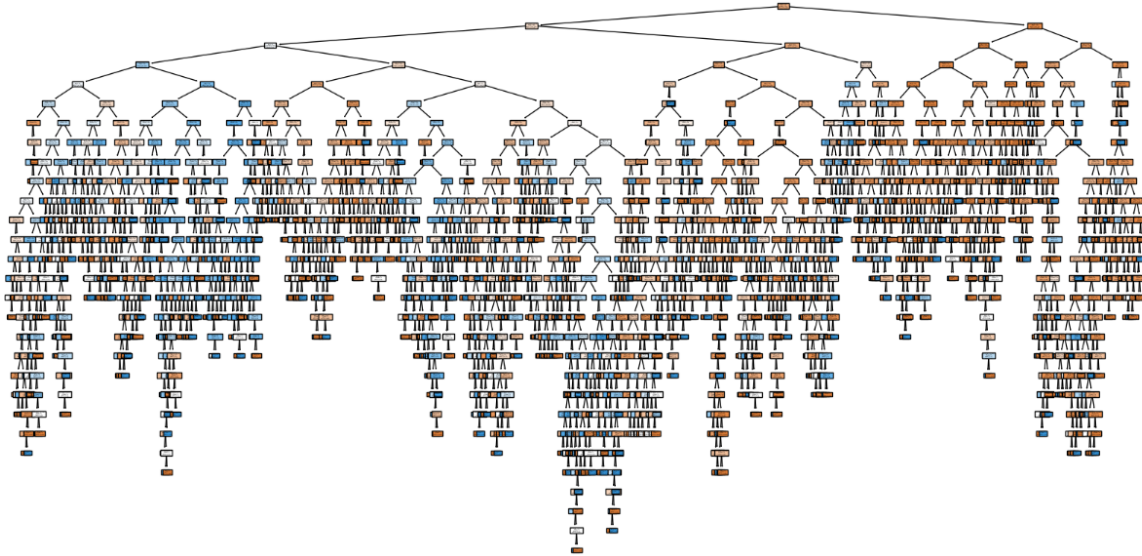
print(f'Accuracy: {accuracy:.2f}')
print(f'Precision: {precision:.2f}')
print(f'Recall: {recall:.2f}')
print(f'F1 Score: {f1:.2f}')

Accuracy: 0.73
Precision: 0.48
Recall: 0.46
F1 Score: 0.47
```

CHECKING THE MODEL AND HOW EFFECTIVE IT IS

- **Accuracy: 0.73:** The model correctly predicts 73% of the total instances.
- **Precision: 0.48:** Of all the positive predictions, 48% are correct.
- **Recall: 0.46:** The model identifies 46% of the actual positive cases.

- **F1 Score: 0.47:** The model's balance between precision and recall is moderate, indicating a need for improvement in both.



This is the visual representation of the decision tree and it is hard to draw conclusions of effects of different parameters from this. Therefore tried using random forests and logistic regression to understand the effect and importance of the other parameters.

RANDOM FOREST MODEL AND LOGISTIC REGRESSION

```

# Random Forest Model
rf_clf = RandomForestClassifier(random_state=42)
rf_clf.fit(X_train, y_train)
y_pred_rf = rf_clf.predict(X_test)
accuracy_rf = accuracy_score(y_test, y_pred_rf)
precision_rf = precision_score(y_test, y_pred_rf)
recall_rf = recall_score(y_test, y_pred_rf)
f1_rf = f1_score(y_test, y_pred_rf)
print(f"Random Forest- \nAccuracy: {accuracy_rf:.2f}\nPrecision: {precision_rf:.2f}\nRecall: {recall_rf:.2f}\nF1 Score: {f1_rf:.2f}")

# importances of different features in the data
importances = rf_clf.feature_importances_
feature_names = X_train.columns
importance_data = pd.DataFrame({'Feature': feature_names, 'Importance': importances}).sort_values(by='Importance', ascending=False)

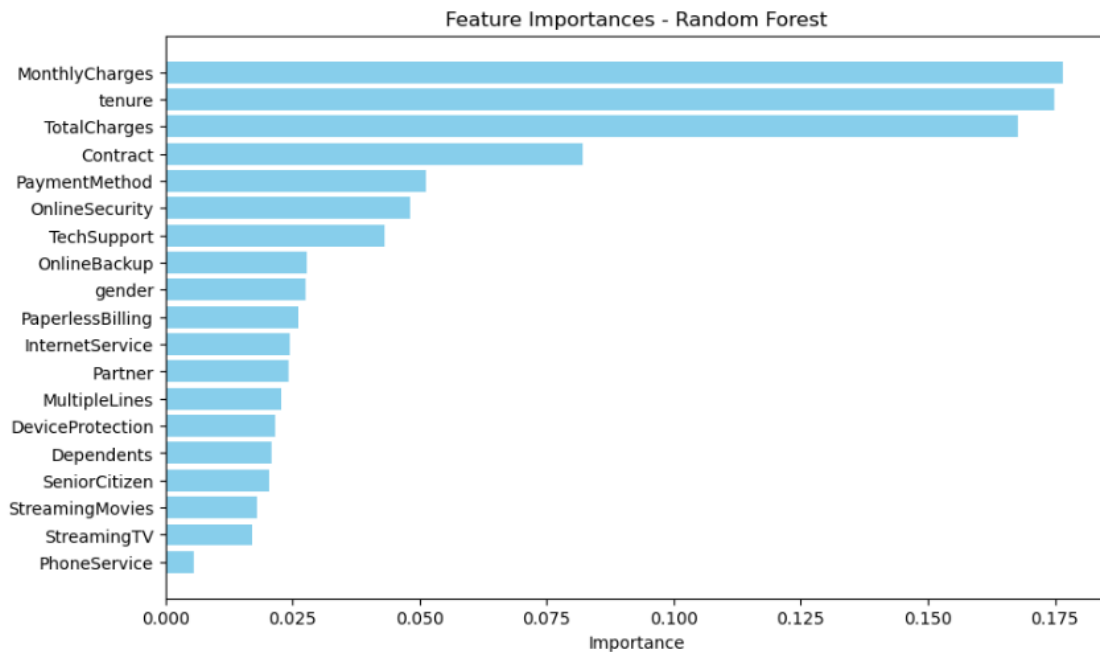
# Plot feature importances
plt.figure(figsize=(10, 6))
plt.barh(importance_data['Feature'], importance_data['Importance'], color='skyblue')
plt.xlabel('Importance')
plt.title('Feature Importances - Random Forest')
plt.gca().invert_yaxis() # Invert y-axis to have the most important feature at the top
plt.show()

# Logistic Regression Model
lr_clf = LogisticRegression(max_iter=1000, random_state=42)
lr_clf.fit(X_train, y_train)
y_pred_lr = lr_clf.predict(X_test)
accuracy_lr = accuracy_score(y_test, y_pred_lr)
precision_lr = precision_score(y_test, y_pred_lr)
recall_lr = recall_score(y_test, y_pred_lr)
f1_lr = f1_score(y_test, y_pred_lr)
print(f"Logistic Regression - \nAccuracy: {accuracy_lr:.2f} \nPrecision: {precision_lr:.2f}\nRecall: {recall_lr:.2f} \nF1 Score: {f1_lr:.2f}\n")

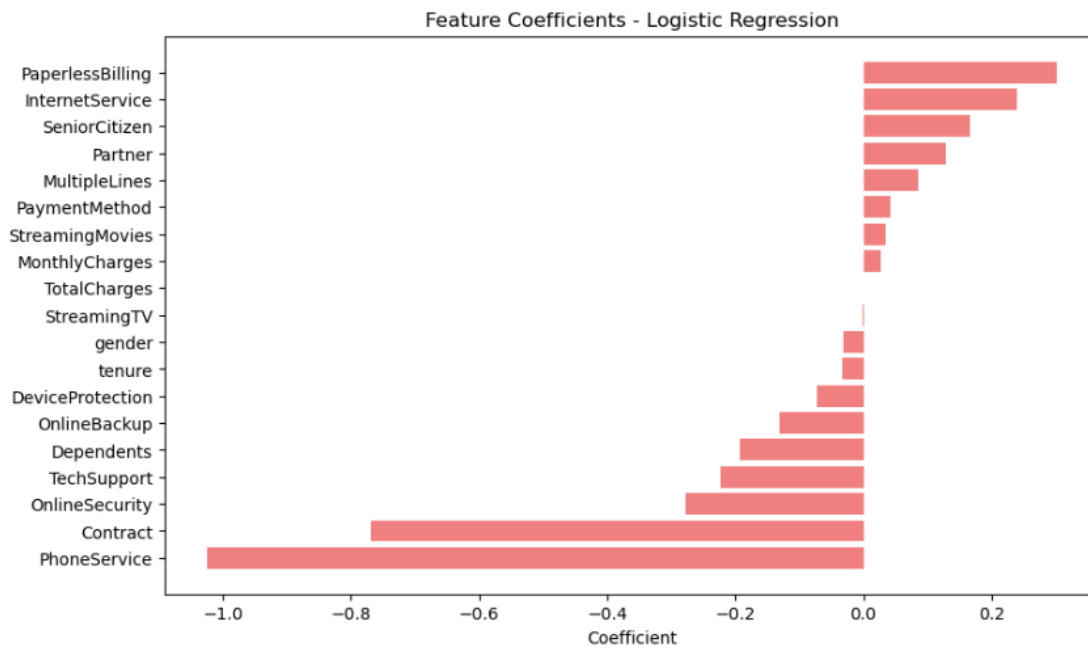
# Coefficients for Logistic Regression
coefficients = pd.DataFrame({'Feature': feature_names, 'Coefficient': lr_clf.coef_[0]}).sort_values(by='Coefficient', ascending=False)

# Plot coefficients
plt.figure(figsize=(10, 6))
plt.barh(coefficients['Feature'], coefficients['Coefficient'], color='lightcoral')
plt.xlabel('Coefficient')
plt.title('Feature Coefficients - Logistic Regression')
plt.gca().invert_yaxis() # Invert y-axis to have the highest coefficient at the top
plt.show()

```



This shows that the monthly charges have the most importance in the purchase decision i.e. churn rate . The importance of features is plotted in decreasing order. This helps us understand which parameters could be dropped while training if the data set is very large.



the coefficient shows the strength of the relationship between feature and target variable.

coefficient of 0.2 indicates - odds ratio = $\exp(0.2)$ which is approximately 1.22 which means there is a slight increase in the odds of churn. if this is negative it indicates with an increase in that particular attribute the churn rate decreases.